



CHALMERS
UNIVERSITY OF TECHNOLOGY

TIF345 / FYM345 Advanced Simulation and Machine Learning

Paul Erhart

Andreas Ekström

Bernhard Mehlig

Arkady Gonoskov



<https://tif345.materialsmodeling.org>

1

Plan for the next two weeks

- 1-2 Foundations: Linear models, Gaussian processes
- 3-4 Applications: Linear models, Gaussian processes, Regression
- 5-7 Neural networks and other ML techniques

L1: Cluster expansions (CEs) and using the covariance matrix

L2: Feature selection algorithms

→ P2a: Regression with CEs

L3: Sensitivity analysis

L4: Global optimization

→ P2b: GPs and Bayesian optimization

tif/fim{345}

Advanced Simulation and Machine Learning

- 5. Advanced regression
- 6. Toward applications
 - 6.1. Alloy cluster expansions
 - 6.2. Phonons and force constants
 - 6.3. Interatomic potentials
- 7. Compressive sensing
- 8. Sensitivity analysis
- 9. Global optimization

CHALMERS
UNIVERSITY OF TECHNOLOGY

Demos

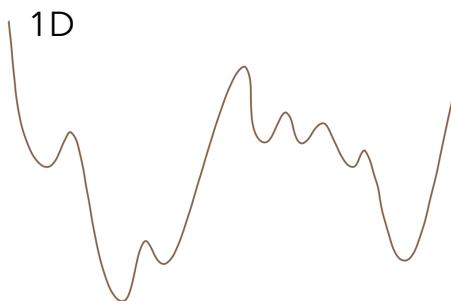


Erhart

Global optimization

3

Global optimization



Stochastic methods



(Meta)heuristic methods

Possible targets

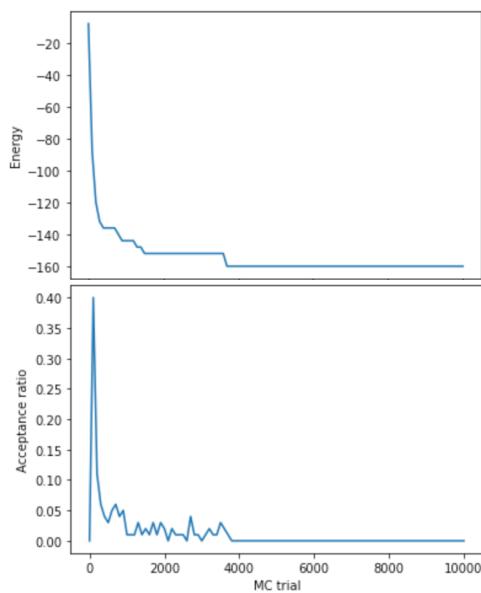
Atomic structures
Model loss functions
Instrument settings
Path finding
Network optimization
...

Response surface methodology-based

Stochastic methods: Simulated annealing

$T_{\text{initial}} \rightarrow T_{\text{final}}$: What is optimal?

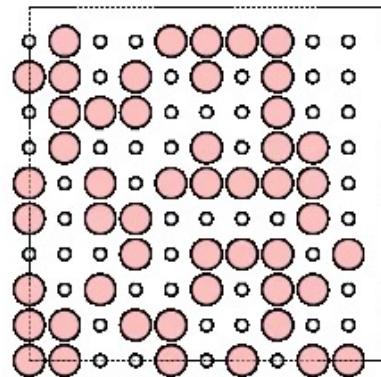
$$T_i = \frac{T_0}{\log(i + i_0)}, i \geq 1$$



$$S = k_B \ln \Omega$$

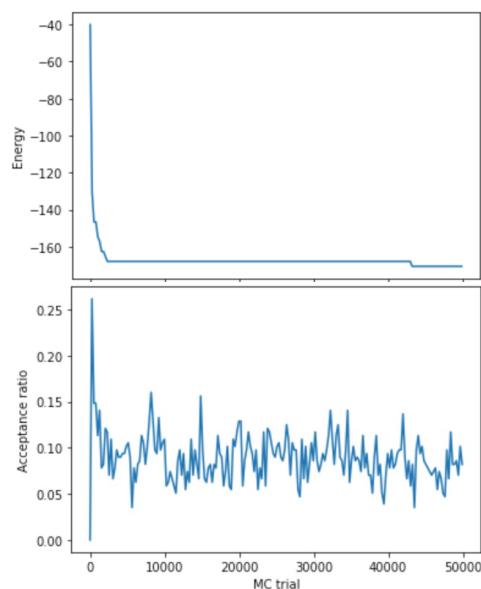
$$F = U - TS$$

$$\mathcal{Z} = \sum_i \exp [-E_i/k_B T]$$



Stochastic methods: Simulated annealing

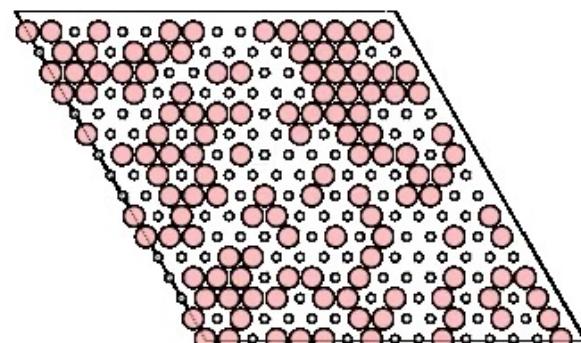
What causes the difference in behavior?



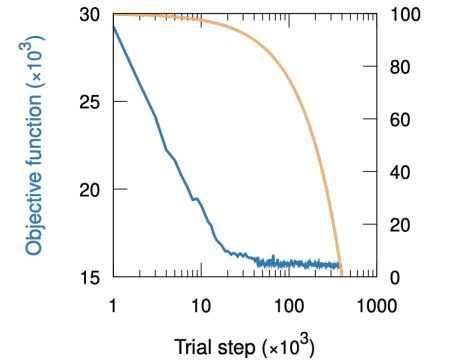
$$S = k_B \ln \Omega$$

$$F = U - TS$$

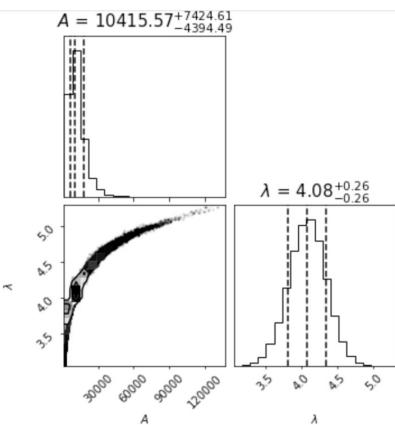
$$\mathcal{Z} = \sum_i \exp [-E_i/k_B T]$$



Stochastic methods: Simulated annealing

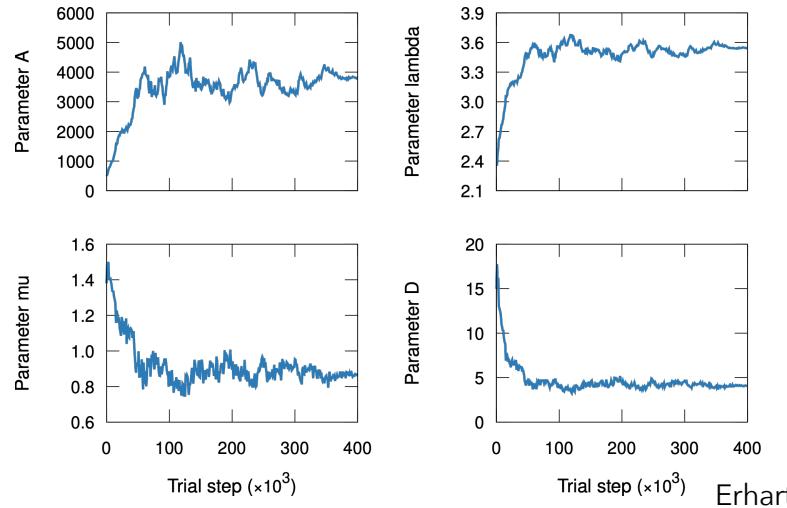


Optimizing a sloppy model



CHALMERS
UNIVERSITY OF TECHNOLOGY

7



Stochastic methods: Critical slowing down

Optimization is probing a *free energy* landscape

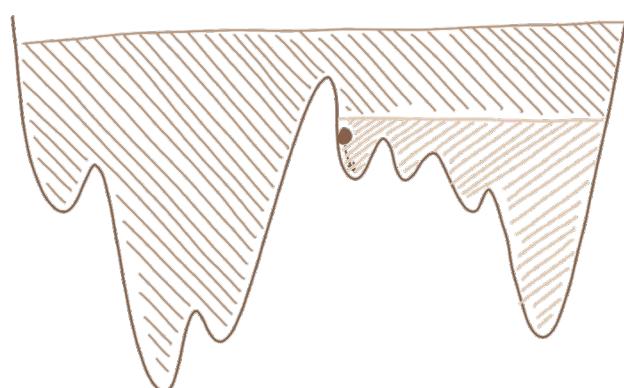
$$S = k_B \ln \Omega$$

$$F = U - TS$$

Zero temperature

High temperature

Lower temperature



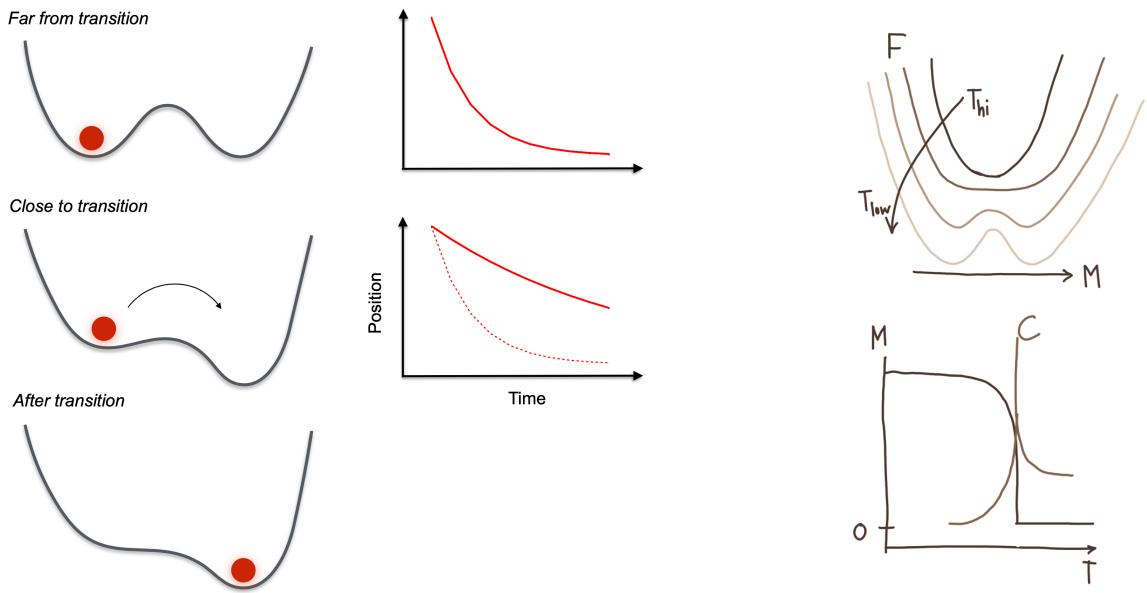
CHALMERS
UNIVERSITY OF TECHNOLOGY

Erhart

o

Stochastic methods: Critical slowing down

Near phase transitions systems recover more slowly from small perturbations



Stochastic methods: Parallel tempering

Parallel tempering (aka Replica exchange) → overcome critical slowdown

N statistically independent replicas
Sampled at temperatures T_i (schedule)
Exchange replicas with probability (satisfying detailed balance)

$$\mathcal{P} = \min \left[1, \frac{\exp \left(-\frac{E_j}{k_B T_i} - \frac{E_i}{k_B T_j} \right)}{\exp \left(-\frac{E_i}{k_B T_j} - \frac{E_j}{k_B T_i} \right)} \right] = \min \left[1, \exp \left((E_i - E_j) \left(\frac{1}{k_B T_i} - \frac{1}{k_B T_j} \right) \right) \right]$$

Acceptance rate at lower T where $R=T'/T$

$$Ac(T, T') \simeq \frac{2}{B(d/2, d/2)} \int_0^{1/(1+R)} \theta^{d/2-1} (1-\theta)^{d/2-1} d\theta.$$

How do you choose R ?

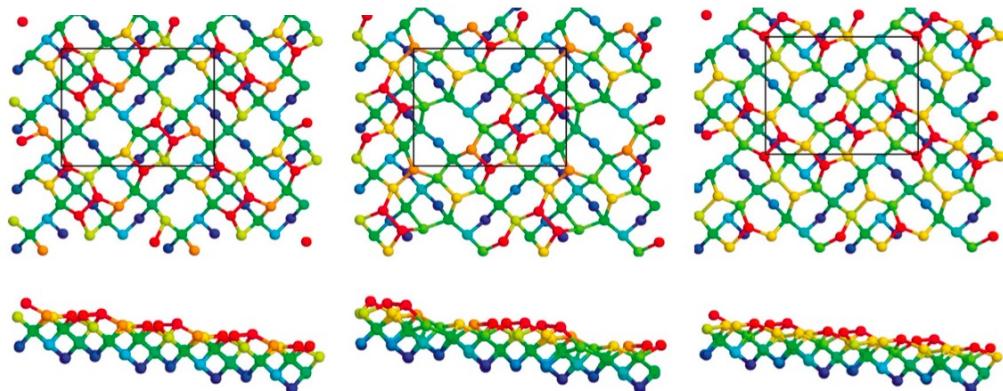
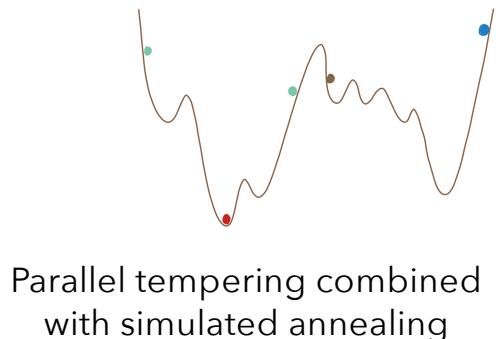
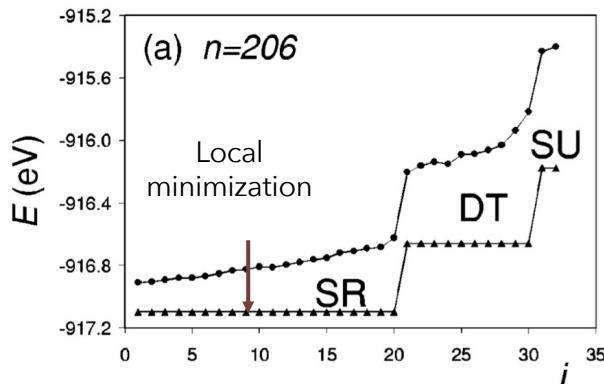
Now: geometric progression

$$T_i = R^{i-1} T_{\min}, 1 \leq i \leq N$$

In simulated annealing (log)

$$T_i = \frac{T_0}{\log(i + i_0)}, i \geq 1$$

Stochastic methods: Example Si (105) surface



CHALMERS
UNIVERSITY OF TECHNOLOGY

Mobano and Predescu, Phys. Rev. B 70, 085211 (2004)

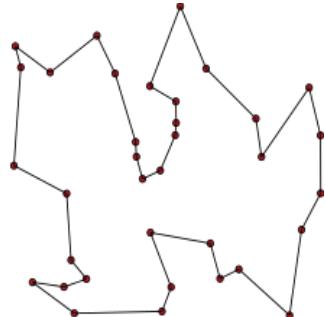
11

(Meta)heuristic methods

1. Particle swarm optimization



2. Ant colony optimization



3. Evolutionary algorithms



CHALMERS
UNIVERSITY OF TECHNOLOGY

Erhart

Particle swarm optimization



https://www.youtube.com/watch?v=Y-5ffI5_7AI

<https://www.youtube.com/watch?v=eakKfY5aHmY> Erhart

13

Particle swarm optimization: Example algo 1

Consider a swarm of particles, each of which

- has a position and a velocity,
- knows the value of the objective function for its position,
- remembers its best previous position
- knows the best positions and respective functions of its neighbors.

$$\begin{aligned}\vec{v}_{i,t+1} &= c_1 \vec{v}_{i,t} + c_2 (\vec{x}_{i-\text{nbr},t} - \vec{x}_{i,t}) + c_3 (\vec{x}_{i-\text{nbr},t} - \vec{x}_{i,t}) \\ \vec{x}_{i,t+1} &= \vec{x}_{i,t} + \vec{v}_{i,t+1},\end{aligned}$$

c_1, c_2, c_3 : social/cognitive confidence coefficients

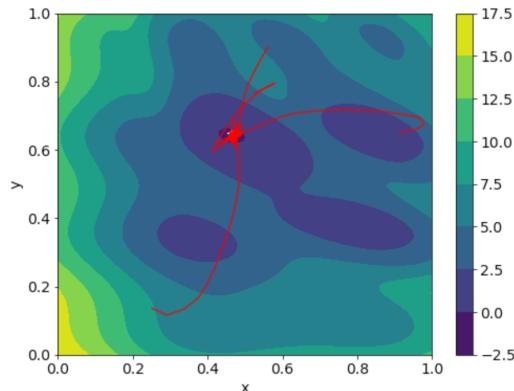
- c_1 quantifies how much the particle trusts *itself* now
- c_2 measures how much the particle trusts its experience
- c_3 determines how much the particle trusts its *neighbors*

Particle swarm optimization: Example algo 2

Consider a swarm of particles, each of which

1. Evaluate the objective function, yielding the value u_j for particle j
 - If u_j is lower than the previous lowest value for this particle, at position \mathbf{p}_j , overwrite \mathbf{p}_j
 - If u_j is lower than the previous lowest value for all particles, at position \mathbf{g} , overwrite \mathbf{g} (global best position)
2. Update the velocity, \mathbf{v}_{new}
 $\mathbf{v}_{\text{new}} = \omega \mathbf{v} + c_1 R_1 (\mathbf{p}_j - \mathbf{x}) + c_2 R_2 (\mathbf{g} - \mathbf{x}),$
3. Update the position, $\mathbf{x}_{\text{new}} = \mathbf{x} + \mathbf{v}_{\text{new}} \Delta t$
 - $c_1 \rightarrow \infty$: exploration
 - $c_2 \rightarrow \infty$: exploitation
 - ω, R_1, R_2 : random numbers

Demo



CHALMERS
UNIVERSITY OF TECHNOLOGY

Images from <https://wikipedia.org>

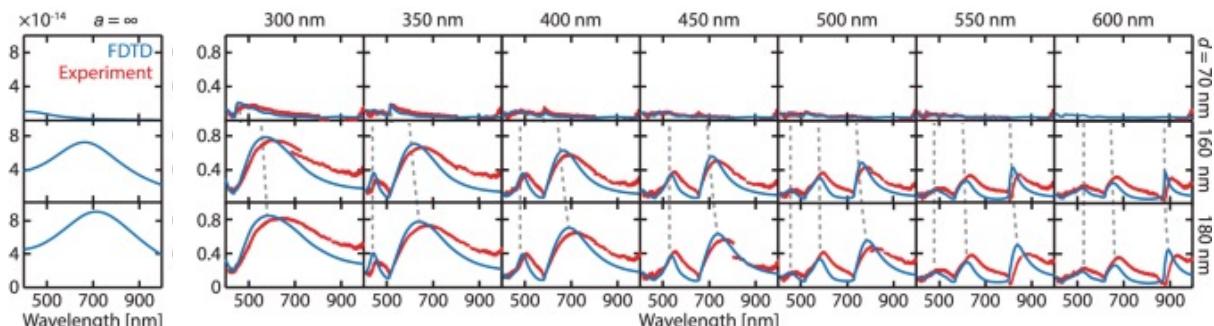
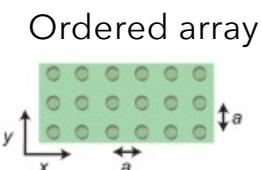
Erhart

15

PSO for optimizing hydrogen sensors

Pd nanoparticles for H₂ detection

Random array

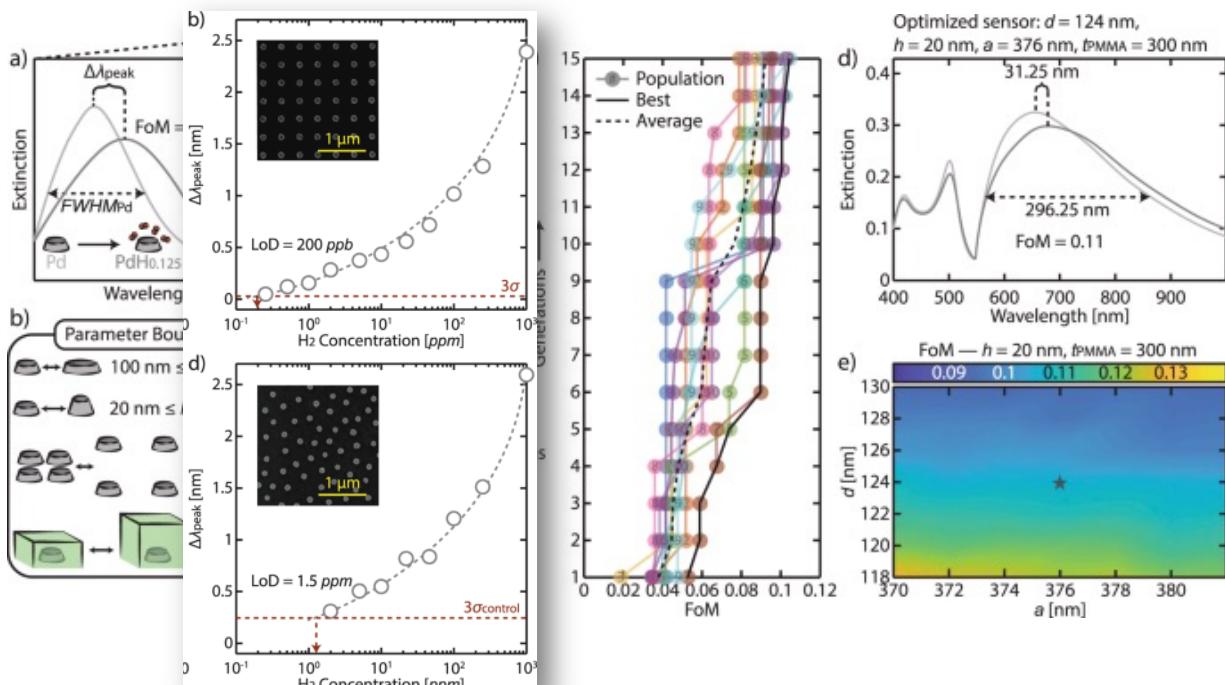


CHALMERS
UNIVERSITY OF TECHNOLOGY

Nugroho et al., Nature Comm. 13, 5737 (2022)

Erhart

PSO for designing nanoparticle arrays



Final result (experimental verification!)

Natural evolution strategies (NES)

Evolutionary algorithms optimize by “fitness”
 Natural ESs operate on multivariate Gaussian mutation distribution
 → Only function (no gradient) evaluations still comparable to 2nd order

What does it do?

It evolves a “search distribution”
 by following a “sampled natural gradient”

- Assume each parameter is given by a (normal) distribution
- The “search distribution” is defined by $\theta = (\boldsymbol{\mu}, \boldsymbol{\Sigma})$ with mean $\boldsymbol{\mu}$ and covariance $\boldsymbol{\Sigma}$
- The gradient of the expected “fitness” can be sampled using the “log-likelihood” trick

The log-likelihood trick

Density of normal search distribution $\mathcal{N}(\mu, \Sigma) \rightarrow \mathcal{N}(\mathbf{0}, \mathbf{1})$

$$p(x | \theta) = \frac{1}{(\sqrt{2\pi})^d \det(A)} \cdot \exp\left(-\frac{1}{2} \|A^{-1} \cdot (x - \mu)\|^2\right) \text{ with } \mathbf{A}\mathbf{A}^T = \Sigma$$

Expected fitness under search distribution [$f(x)$ = fitness measure]

$$J(\theta) = \mathbb{E}[f(x) | \theta] = \int f(x) p(x | \theta) dx$$

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \int f(x) p(x | \theta) dx$$

$$= \int f(x) \nabla_{\theta} p(x | \theta) dx$$

$$= \int f(x) \nabla_{\theta} p(x | \theta) \frac{p(x | \theta)}{p(x | \theta)} dx \quad \begin{matrix} \text{Sample via MC (derivative} \\ \text{can be computed efficiently)} \end{matrix}$$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{n} \sum_{i=1}^n f(x_i) \nabla_{\theta} \log(p(x | \theta))$$

Further improvements

Use “natural” gradient

$$G = F^{-1} \nabla_{\theta} J(\theta) \quad F: \text{Fisher information metric}$$

Fitness shaping: fitness \rightarrow utility

$$\nabla_{\theta} J(\theta) \approx \frac{1}{n} \sum_{i=1}^n f(x_i) \nabla_{\theta} \log(p(x | \theta))$$

$$\Rightarrow \nabla_{\theta} \widehat{J}(\theta) = \sum_{i=1}^n u_i \nabla_{\theta} \log(p(x_i | \theta))$$

$u_1 \geq \dots \geq u_n$: ranked utility values for each member of population

Further improvements

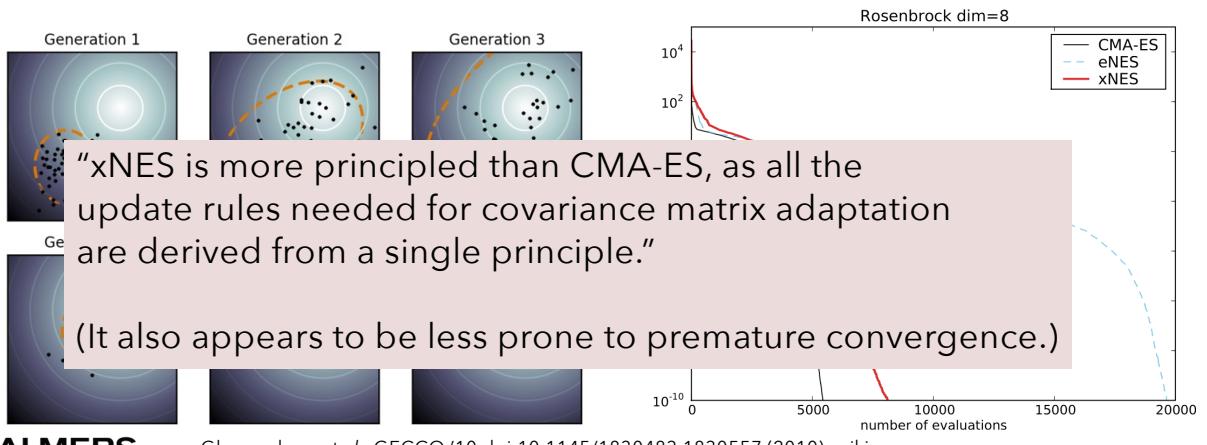
eNES = exact NES

- Improved handling of Fisher information metric
- Introduces important mixing

xNES = exponential NES

- Exponential space instead of directly on Σ to ensure positive-definite Σ)
- Avoids computation of Fisher metric by transforming coordinate system

Covariance matrix adaptation evolution strategy (CMA-ES)



22

xNES by example

Task: Training a neural network potential (beyond SGD)

$$N_{\text{par}} = (N_{\text{des}} + 2)N_{\text{neu}} + 1 \quad \sim \text{typically a few thousand}$$

N_{des} : Number of nodes in input (descriptor) layer

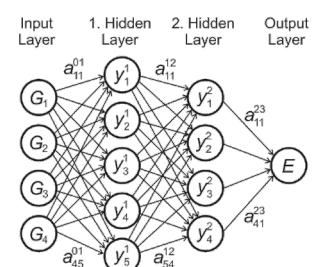
N_{neu} : Number of neurons in hidden layer

Loss function

$$L(\mathbf{z}) = \lambda_e L_e(\mathbf{z}) + \lambda_f L_f(\mathbf{z}) + \lambda_v L_v(\mathbf{z}) + \lambda_1 L_1(\mathbf{z}) + \lambda_2 L_2(\mathbf{z})$$

energies forces virials ℓ_1 ℓ_2

\mathbf{z} : parameter vector



xNES by example

1. Initialization

2. Loop over generations

1. Create a population of solutions \mathbf{z}_k ($1 \leq k \leq N_{\text{pop}}$)

$$\mathbf{z}_k \leftarrow \mathbf{m} + \mathbf{s} \odot \mathbf{r}_k \quad \mathbf{r}_k \sim \mathcal{N}(0, 1)$$

2. Evaluate loss function and sort populations from small to large

- ### 3. Update natural gradients

$$\text{Update natural gradients} \quad (\boldsymbol{u}_i \text{ rank-based utility values}) \quad \nabla_{\mathbf{m}} J \leftarrow \sum_{k=1}^{N_{\text{pop}}} \boldsymbol{u}_k \mathbf{r}_k,$$

- #### 4. Update mean and sdv

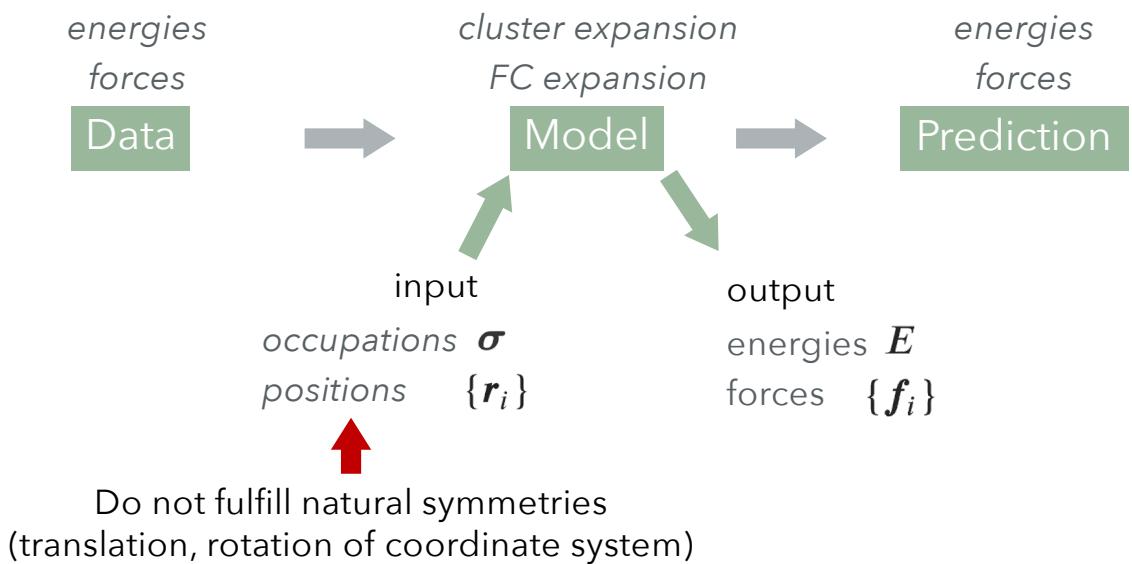
$$\mathbf{m} \leftarrow \mathbf{m} + \eta_{\mathbf{m}} (\mathbf{s} \odot \nabla_{\mathbf{m}} J) \quad \nabla_{\mathbf{s}} J \leftarrow \sum_{k=1}^{\text{pop}} u_k (\mathbf{r}_k \odot \mathbf{r}_k - 1)$$

$$\mathbf{s} \leftarrow \mathbf{s} \odot \exp\left(\frac{\eta_s}{2}\nabla_{\mathbf{s}}J\right), \quad \eta_s, \eta_m: \text{"learning rates"}$$

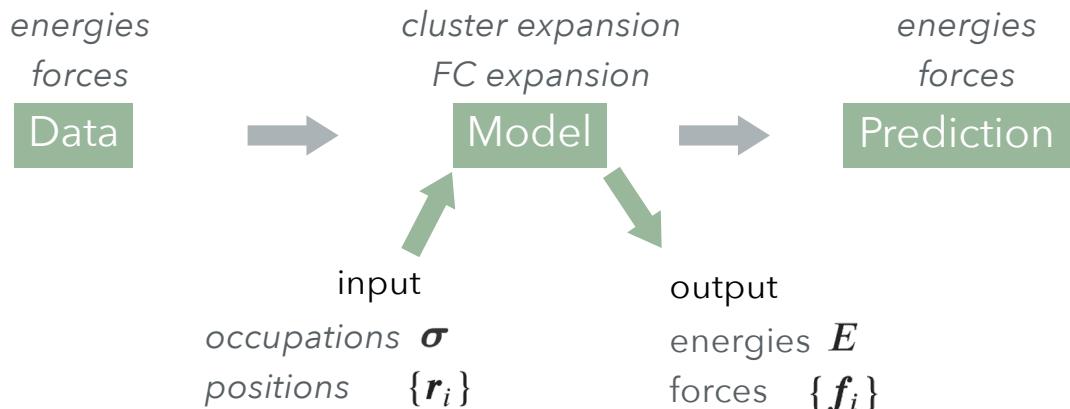
24

Injection: Machine learned interatomic potentials

How do we represent the state of a system?



How do we represent the state of a system?



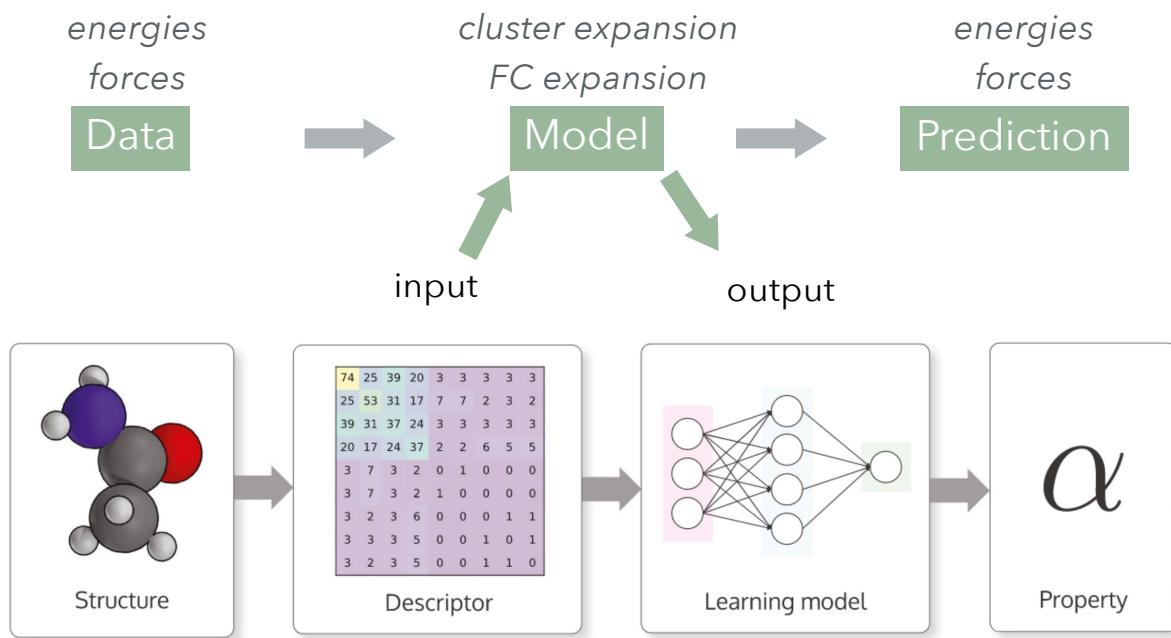
Consider a cluster expansion

$$E(\sigma) = E_0 + \sum_i J_i \sigma_i + \sum_{ij} J_{ij} \sigma_i \sigma_j + \sum_{ijk} J_{ijk} \sigma_i \sigma_j \sigma_k + \dots$$

$$E_{CE} = J_0 + \sum_{\alpha} m_{\alpha} J_{\alpha} \Pi_{\alpha}(\sigma)$$

Coordinate transformation

How do we represent the state of a system?



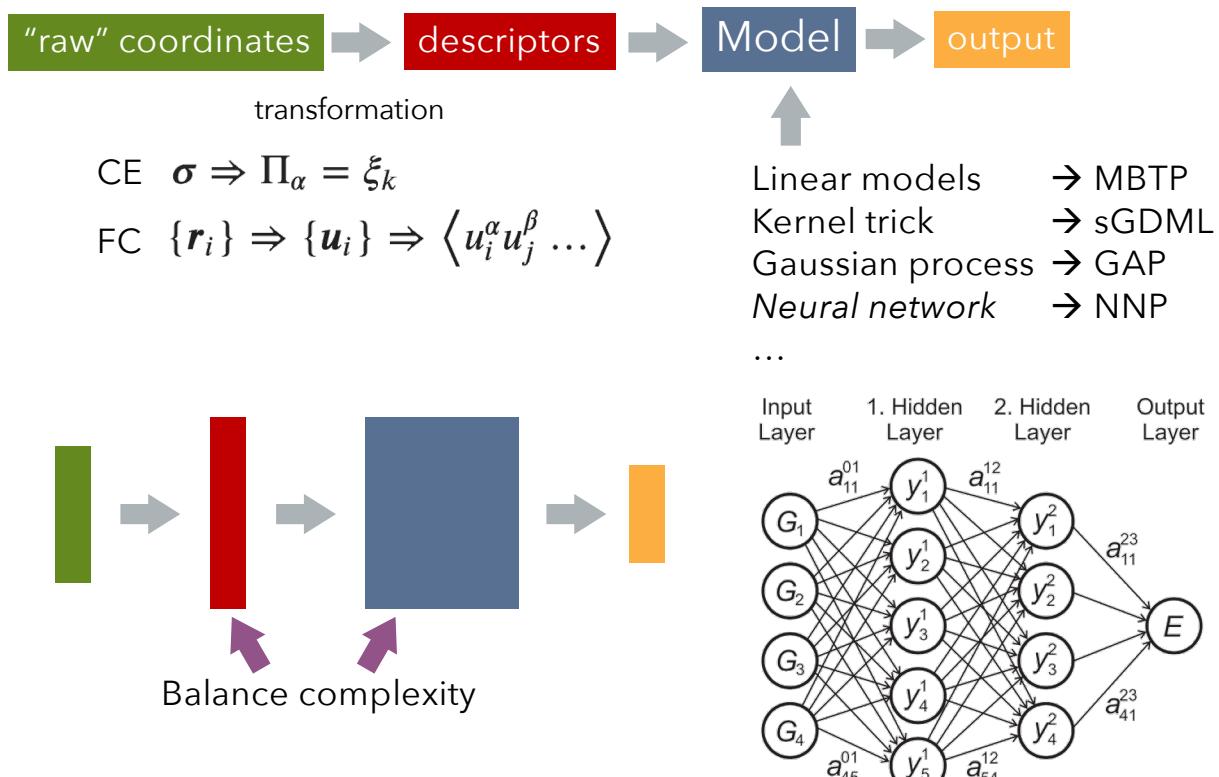
CHALMERS
UNIVERSITY OF TECHNOLOGY

Figure from Comp. Phys. Comm. 247, 106949 (2020)

Erhart

28

Model vs descriptor complexity



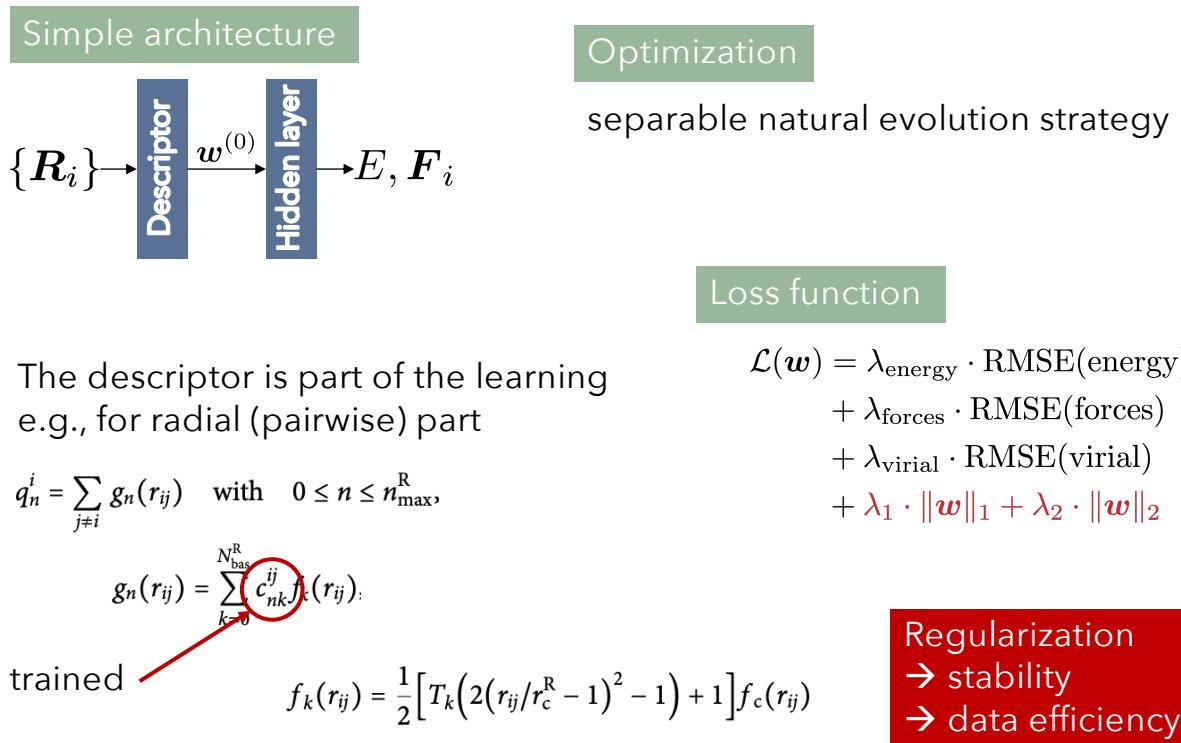
CHALMERS
UNIVERSITY OF TECHNOLOGY

Figure from J. Chem. Phys. 134, 074106 (2011)

Erhart

29

Neuroevolution potential (NEP) models



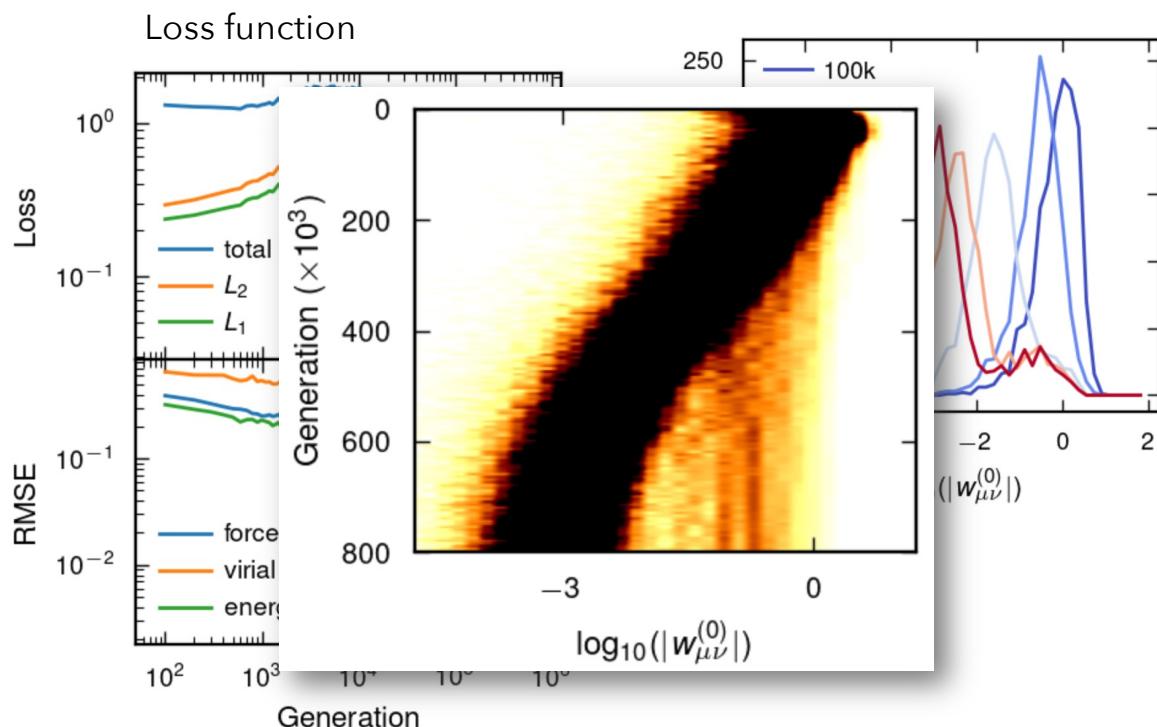
CHALMERS
UNIVERSITY OF TECHNOLOGY

Fan et al., J. Chem. Phys. 157, 114801 (2022)

Erhart

32

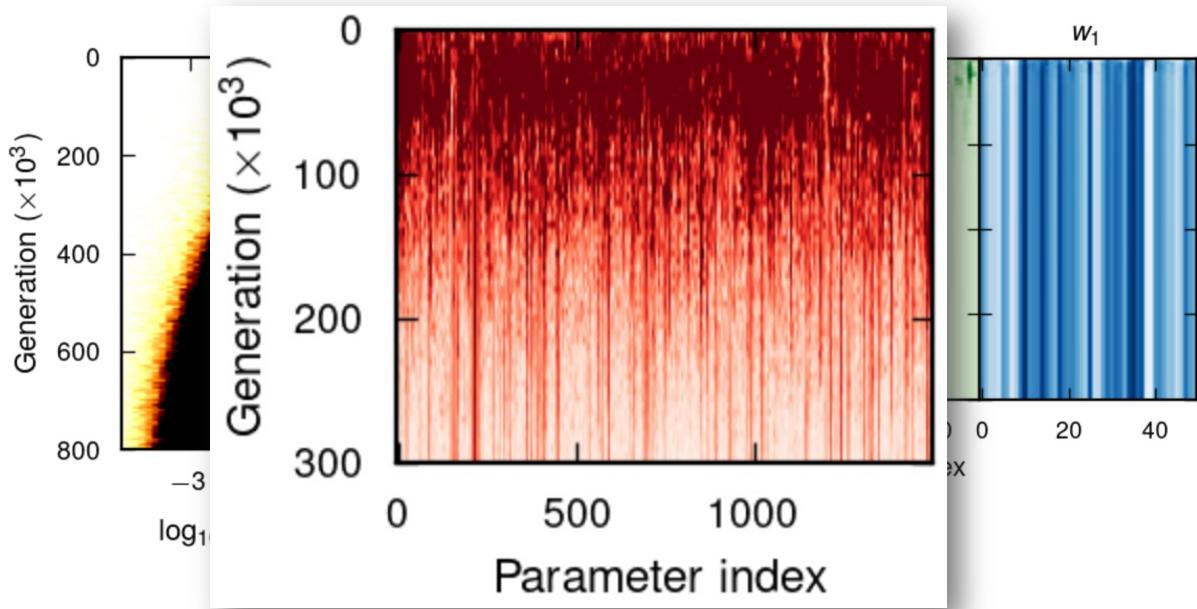
Evolution of sparsity during training



CHALMERS
UNIVERSITY OF TECHNOLOGY

Model for $\text{Cs}_{1-y}\text{Rb}_y\text{PbCl}_{3(1-y-z)}\text{Br}_{3y}\text{I}_{3z}$ Erhart

Feature selection

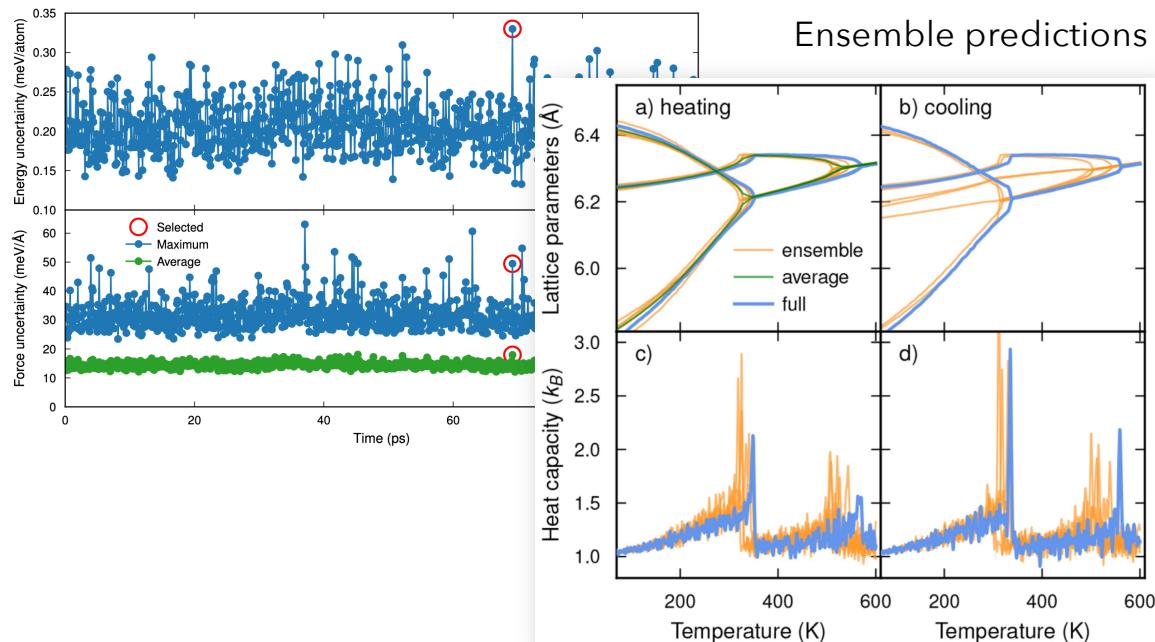


Note the “jumps” between solutions

In the final model only 25% of the NN weights are non-zero

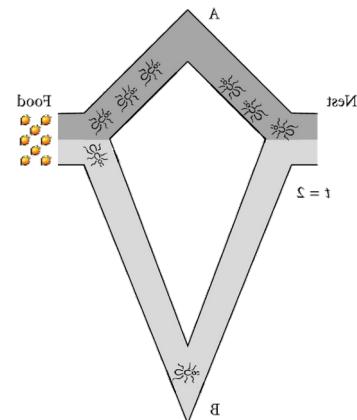
Model ensemble through k-fold

Active learning



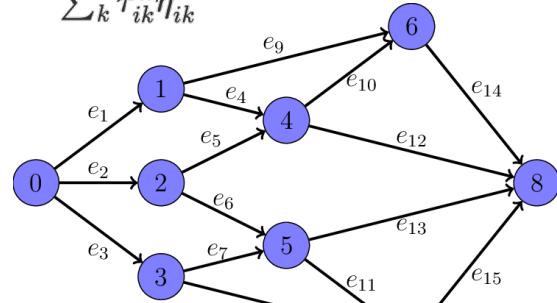
Ant colony optimization

1. Walkers are placed at origin and stochastically find pathways through graph
2. Quality (weighted length) of paths is calculated
3. "Pheromone" is deposited along each path, with more pheromone placed on shorter paths
4. Next set of walkers chooses each next step weighing in both the desirability (typically related to the length) of the path and the amount of pheromone already deposited



Ant colony optimization

Probability for choosing a step $\mathcal{P}_{ij} = \frac{\tau_{ij}^\alpha \eta_{ij}^\beta}{\sum_k \tau_{ik}^\alpha \eta_{ik}^\beta}$



Pheromone update

$$\tau_{ij} \leftarrow (1 - \rho)\tau_{ij} + \sum_k \Delta\tau_{ik}^k$$

$$\Delta\tau_{ik}^m = \begin{cases} \frac{Q}{L_m} & \text{if ant } m \text{ used edge } (i, k) \text{ during its current tour} \\ 0 & \text{otherwise} \end{cases}$$

Demo 1

α, β, Q are adjustable constants

Demo 2

Response surface methodology-based

Replace complex objective with (very) cheap emulator
→ optimize using emulator via optimizer or simply grid search

Bayesian optimization: use priors over unknown functions
→ Gaussian processes or Kriging

Algorithm

1. find the next sampling point by optimizing the acquisition function
2. sample the objective function
3. add the sample to the data and update the GP
4. return to 1.

Choices

- dimensionality of space (descriptor, reaction coordinate, ...)
- acquisition function/strategy

$$\mathbf{x}_{\text{new}} = \arg \min_{\mathbf{x}} [f(\mathbf{x}) + \beta \text{var}_f(\mathbf{x})] \quad \begin{array}{l} \beta \rightarrow 0: \text{exploitation} \\ \beta \rightarrow \infty: \text{exploration} \end{array}$$

Summary



Stochastic methods

Simulated annealing
Replica exchange



(Meta)heuristic methods

Evolutionary methods (including genetic algorithms)
Ant colony optimization
Particle swarm optimization

Exploitation vs exploration

Possible targets

Atomic structures
Model loss functions
Instrument settings
Path finding
Network optimization
...

Response surface methodology-based

Bayesian optimization